



Accuracy Gains Without Training Your Own Model: Parallel Reasoning for LLM Apps

quick intro 🖐️

Eugene Cho

Member of Technical Staff

- I work on the core post-training team
- Prev worked a lot on synthetic data + reasoning



some helpful pre-requisite knowledge

- a large language model (LLM) auto-regressively predicts the next token given some past text
 - for simplicity's sake, you can think of a token as a word
 - modern day LLMs are specifically trained to act like agentic assistants to user queries (done through post-training)
- LLMs are NON deterministic
 - that is, if you ask an LLM the exact same question multiple times, chances are the output will be different
 - the degree of non-determinism/randomness is steerable through a parameter called **temperature**. Higher temperature == more random

Agenda

01

what is "test-time compute"?

02

scaling test-time compute

03

verifiers & how to implement them for your usecase

04

some demos
(including the tradeoffs)



01

what is “test-time compute”?

“train-time compute”

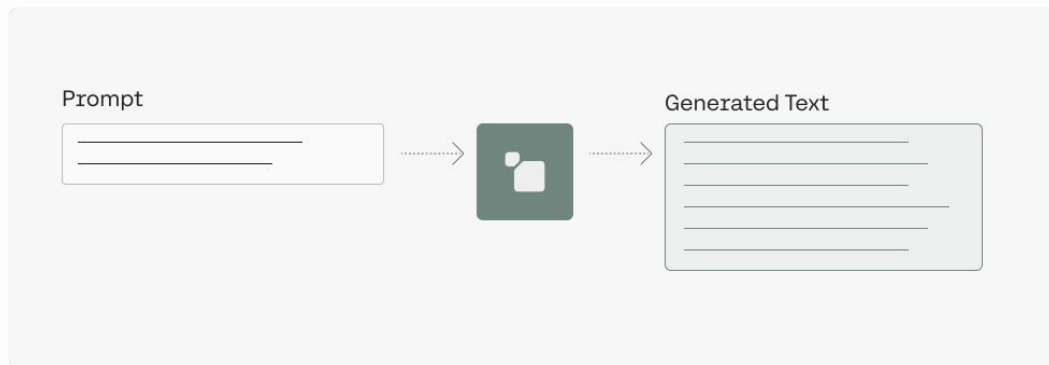
The amount of computational power used to train an AI model (i.e. compute gradients, update weights, etc.)

“test-time compute”

The amount of computational power used by an AI model at “inference” or “test”-time (i.e. using a model to provide predictions on new data)

in our case

test-time compute in LLM applications refers to the amount of energy spent generating response(s) to a user request/prompt, or more generally, some kind of agentic (tool calling) workflow





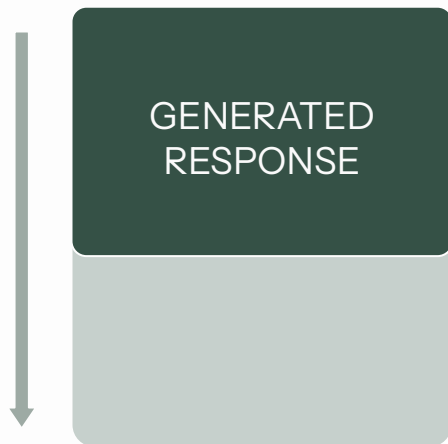
02

scaling test-time compute

two ways to scale

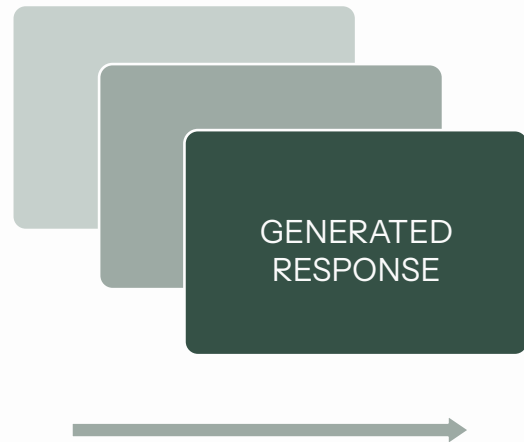
vertically

increasing response
trajectory length



horizontally


increasing the number of
response trajectories



vertical scaling of ttc (brief)

there are a few ways to do this:

1. prompt the model to think step by step ([Chain of Thought Prompting](#))



2. use a reasoning model (trained with RL which learns to natively reason for longer) ([DeepSeek R1 Paper](#))
 - a. most reasoning models allow you to control the amount of thinking it does via a thinking budget/effort API param.





Figure 1. Test-time scaling with s1-32B. We benchmark s1-32B on reasoning-intensive tasks and vary test-time compute.


- Literature has established scaling laws that as we increase response generation length, **accuracy/performance** increases!

horizontal scaling of ttc (focus of today)

we can also just generate multiple responses for the same prompt. using these responses, we can then design a method of selecting or creating the best response based on some definition of correctness to help boost accuracy

There are a few names / variations of this:

- “best-of-N”
- parallel reasoning
- maj@k (self-consistency)



why does this work?

- recall LLMs are **non-deterministic**, so it could be the case that *sometimes* the model will answer a prompt “correctly” and *sometimes* it won’t
 - For many tasks, a well aligned LLM will not require many “trials” to get at least one correct trajectory

- for certain use cases, verifying correctness is often an easier task than answering the question correctly



03

verifiers

what is a verifier?

for the purpose of this talk, a **verifier** is an entity that assigns a label of correctness to a single LLM trajectory.

- think of it as an arbitrary function v which takes in an LLM trajectory t as input, and spits out a boolean (True or False), representing the correctness of t . That is:

$$v: t \rightarrow \{\text{True}, \text{False}\}$$

- Note: in practical and more complex cases, the verifier function can output any scalar



implementing my verifier

it all boils down to defining “correctness” in your use case

- Can be as simple as:
 - A property of the LLM response meeting some criteria
 - Provides code that passes a pre-defined unit test
 - Response contains a specific desired keyword
 - A change in some external state (will come up if model has access to tools)
 - A email is sent successfully
 - A status is toggled off
- Can be as complex as:
 - A combination of simple conditions like the above:
 - I.e. **only if** all the conditions above are met, then the response is correct
 - Using an LLM to provide a label based on a more subjective criteria
 - ex: let's say my usecase is a chatbot that outputs encouraging messages, I can use another LLM to verify whether or not the message is encouraging or not!

synthesis: how to select the “best” response?

- it's use case dependent
 - you'll likely need to get your hands dirty and try a few different applicable approaches
- some popular synthesization approaches (easiest to hardest):
 - **Ground-truth programmatic verification**
 - Use this when it is easy to programmatically verify the correctness of an **arbitrary response** (ex. Basic textual properties, quantitative measurements)
 - Here your verifier directly selects your best response !
 - **take the majority answer (maj@k)**
 - Use this when the correctness of outcomes/answers are not known ahead of time, but the space of possible outputs is objective and atomic such that comparing amongst them is simple (ex. Integer answers)
 - Also is only practical if you suspect the model gets the answer right *most* of the time
 - **ask another LLM (or a *panel* of LLMs) to choose the best one (most complex and hand-wavy)**
 - Use this when outcomes are more difficult to objectively or programmatically compare (ex. writing quality)

best ML practices

BEFORE DOING ANYTHING:

- Define your verifier
- Define an evaluation set
 - this consists of prompts + some definition of a ground truth / correct outcome
 - the larger the evaluation set the better, but in the interest of time you may want to keep it small (< 100 samples)
- evaluate your synthesization methods (including your baseline: which is where you just give the model one chance to solve the problem!)
 - **Evaluate your baseline FIRST before implementing anything else to save yourself time)**
 - If your baseline performance is mid, only then should you consider using test-time scaling!

how do i know if scaling test-time compute will be fruitful?

a checklist

- if there is a clear way to define “correctness” in your usecase ✓
- if the baseline performance of the model i’m using is trash ✓
- if there is some non-zero probability that the LLM can perform the task correctly ✓
 - The higher the probability, the better!

If the checklist is not met, then horizontal scaling will likely **not** be fruitful 💔

- what do i do then?
 - try again with a larger / stronger model (more expensive! 😞)
 - might need to do task specific fine-tuning (doomed for a hackathon 😞)

04

demo (including tradeoffs)

toy problem

Spec:

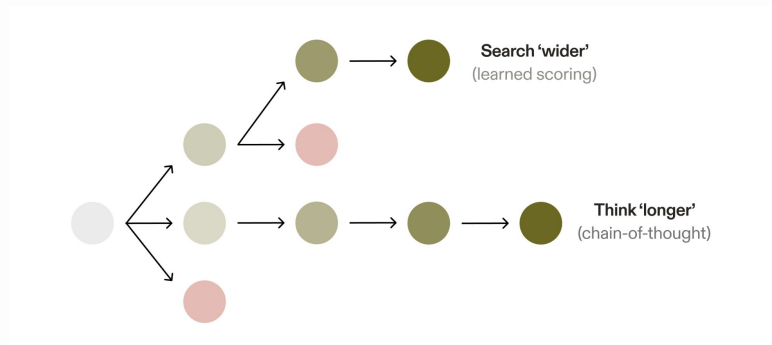
- I have an application that fills in job applications for people.
- The problem:
 - many text boxes in job applications have word count limits.
- The LLM I'm using out of the box is not always good at adhering to the word count limit, despite me asking it to :(!
- For simplicity:
 - We will assume that all the responses the LLM gives are **high quality!** Realistically, this will not be the case, and extra work should be done into prompt engineering, etc.
 - We will also assume that the longer response, the higher quality it is! This is also not really the case!

Follow along here:

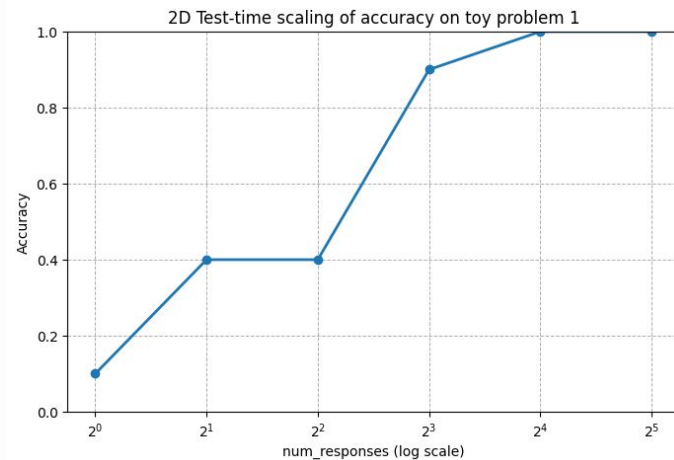
https://github.com/cohere-ai/htn-2025-techtalk/blob/master/toy_problem.ipynb

final food for thought 🤔

What happens if we scale test-time compute both vertically and horizontally at the same time? (2 dimensionally)



even more gains!!! but you are going to go
bankrupt 😞.
but if it is worth it, then go ahead!



An abstract background featuring large, overlapping, rounded shapes in shades of purple, teal, and dark green. The shapes have a soft, blurred appearance, creating a modern and artistic look.

QUESTIONS?

THANKS FOR LISTENING